

FINITE-MEASUREMENT TRANSPORT OF PERTURBATION RELIABILITY RANKINGS

Anonymous authors

Paper under double-blind review

ABSTRACT

Predictions validated in one biological context are often considered for use in another, yet perturbations with low prediction error in one context need not retain their relative accuracy after a context change. We study this ordering of perturbation-level realized error while the predictor is held fixed. Because finite cell counts can themselves induce apparent rank changes, we define a tie-aware discrepancy that compares pairwise ordering distributions and subtracts expected within-context ordering disagreement. The resulting finite-measurement estimand is zero when the two pair-state laws agree, while its unbiased estimator may be negative at finite sample size. In the Frangieh Perturb-CITE-seq dataset (Frangieh et al., 2021), reliability rankings change across control, co-culture and IFN γ conditions under three complementary risk metrics. The changes exceed same-context sampling baselines, become increasingly resolvable with cell budget, and are reproduced in independent HepG2–Jurkat and iPSC–neuron raw-cell analyses (Nadig et al., 2025; Tian et al., 2019). Matched source-only RBF, direct-MLP, latent-MLP and strict source-frozen GEARS sensitivities preserve positive transport signs on their respective eligible Frangieh surfaces, although their within-metric profiles are not identical. Reliability reordering also changes finite top- k experimental shortlists. In the tested controls, however, D_{adj} did not improve held-out regret prediction beyond metric identity and mean-risk shift. A known-truth comparator audit shows why the finite-measurement correction matters: uncorrected rank statistics can declare transport under a same-context null, whereas matched-null corrections restore calibration. The framework therefore supports measurement-aware auditing of cross-context changes in model-error ordering.

1 INTRODUCTION

Single-cell perturbation screens make it possible to measure many interventions in parallel, but exhaustively measuring them across biological contexts is often impractical. Pooled Perturb-seq established a route from genetic perturbations to high-dimensional single-cell readouts (Dixit et al., 2016; Adamson et al., 2016); genome-scale screens and harmonized perturbation resources have since expanded both the intervention universe and the range of measurement depths (Replogle et al., 2022; Peidli et al., 2024). A computational predictor can therefore help determine which source-validated predictions may be reused after a context change.

That narrower reuse decision is an ordering problem. Existing perturbation-prediction studies usually ask whether a model reconstructs an unseen response or generalizes to an unseen perturbation. Representative approaches span latent-variable, optimal-transport, graph-based and foundation-model methods, including scGen, CPA, CellOT, GEARS, scGPT and TxPert (Lotfollahi et al., 2019; 2023; Bunne et al., 2023; Roohani et al., 2024; Cui et al., 2024; Wenkel et al., 2026). Benchmarks broaden this comparison across datasets and metrics (Wei et al., 2026), while direct audits show that complex models need not outperform simple linear baselines under a fixed protocol (Ahlmann-Eltze et al., 2025; Wong et al., 2025). These studies evaluate response fidelity; they do not directly ask whether the identity of reliably predicted perturbations survives a biological context shift.

A ranking shift is not automatically a biological-context effect. With finite numbers of cells, perturbation risks fluctuate, and two observed orders can differ even when the underlying order is not resolved. The number of cells per perturbation–context pseudoreplicate is therefore part of the in-

Fig. 1 Reliability ordering under biological context shift




Figure 1: Reliability ordering under biological context shift. The source predictor is frozen before evaluation in the target context. The rank braid retains all 243 background rank trajectories and highlights a deterministic set of source-rank quantile exemplars for legibility. The compact shortlist preview reports 25 perturbations per context: 6 retained and 19 newly entering the target shortlist. The full 243-label accounting remains in the Fig. 1 source table, while the detailed replacement flow is reserved for Fig. 4A. The decomposition uses the same source-frozen predictor and prespecified evaluation protocol on its declared eligible transfer set. Metric colors are fixed throughout the paper.

ferential problem, not merely technical details (Peidli et al., 2024; Squair et al., 2021). Work on ranking under noisy measurements makes the same point for uncertain and tied orders (Zuk et al., 2007; Sankaran et al., 2025). Our question is complementary to calibration and selective prediction: after accounting for this finite-cell ordering floor, does the relative ordering of realized model risk change between biological contexts?

We answer this question with a source-frozen protocol. Predictions are fixed before target outcomes are observed; context-specific risks are then compared through a tie-aware, measurement-adjusted discrepancy. For each perturbation pair, the estimator compares the distributions of lower, tied and higher risk states across contexts and subtracts expected within-context ordering disagreement using an unbiased order-2 U-statistic (Hoeffding, 1948). The primary analysis uses the Frangieh context surface, same-context pseudoreplicate controls, independent HepG2–Jurkat and iPSC–neuron raw-cell replications, matched-cell-budget resampling and a finite top- k decision layer.

We make three contributions. First, we formalize reliability transport as a finite-measurement comparison of perturbation-pair ordering distributions rather than mean response fidelity. Second, we show that context-associated reordering exceeds a matched sampling floor, has an explicit measurement boundary, and recurs across independent biological systems and matched nonlinear-predictor sensitivities. Third, we quantify its consequence for finite experimental shortlists while keeping post-hoc transport measurement separate from prospective source-only prediction.

2 PROBLEM SETTING AND MEASUREMENT-ADJUSTED ORDERING

Let $X_{p,e}^{(r)}$ denote the raw cells sampled for perturbation p in context e and Monte Carlo replicate r . Each cell is library-size normalized to 10^4 , transformed by $\log(1+x)$, and aggregated into the declared gene-panel response $y_{p,e}^{(r)}$ (a perturbation-versus-control delta for delta-based metrics). The primary predictor is an out-of-fold map: for perturbation p , $f_{s,-\text{fold}(p)}$ is fit only to labels outside p 's fixed fold in the source training context s , using a ten-component PCA representation, ridge penalty $\alpha = 0.1$, and three independently seeded members. It produces a fixed held-out vector $\hat{y}_{p,s}$ before any target outcome is examined. For metric m , the realized risk is

$$R_{p,e}^{(m,r)} = g_m(\hat{y}_{p,s}, y_{p,e}^{(r)}), \quad (1)$$

where larger values mean worse prediction and the direction $s \rightarrow e$ names the training context and evaluation context, respectively. We use one minus cosine similarity (delta cosine), one minus Systema centroid accuracy (Vinas Torne et al., 2026), and one minus absolute-effect rank agreement; complete metric definitions, comparator universes, and tie tolerances are given in Appendix B. These risks are evaluated after target outcomes are observed; source-only uncertainty is kept separate from the ordering estimand.

For a pair $a = (p, q)$, define the tie-aware pair state

$$Z_{a,e}^{(m,r)} = \text{sign}(R_{p,e}^{(m,r)} - R_{q,e}^{(m,r)}) \in \{-1, 0, +1\}. \quad (2)$$

The resampling distribution gives $\pi_{a,e}^{(m)} = [\Pr(Z = -1), \Pr(Z = 0), \Pr(Z = +1)]$. For two contexts e_1, e_2 ,

$$D_{\text{cross}} = \mathbb{E}_{a \sim w}[1 - \pi_{a,e_1}^\top \pi_{a,e_2}], \quad D_{\text{within},e} = \mathbb{E}_{a \sim w}[1 - \|\pi_{a,e}\|_2^2]. \quad (3)$$

The measurement-adjusted discrepancy, equivalently one half of the squared delta-kernel MMD (Gretton et al., 2012), is

$$D_{\text{adj}} = D_{\text{cross}} - \frac{D_{\text{within},e_1} + D_{\text{within},e_2}}{2} = \frac{1}{2} \mathbb{E}_a \|\pi_{a,e_1} - \pi_{a,e_2}\|_2^2 \quad (4)$$

Here w is uniform over the declared unordered perturbation-pair universe and the delta kernel $k(z, z') = \mathbb{I}\{z = z'\}$ compares the three-valued pair states. Thus D_{adj} is the squared separation between the two ordering distributions after removing their within-context coincidence terms. It is a transport quantity for relative reliability, rather than an average change in prediction accuracy. For interpretability, write $t_e = \pi_{a,e}(0)$ and $b_e = \pi_{a,e}(+1) - \pi_{a,e}(-1)$. The pairwise squared separation has the exact decomposition

$$\begin{aligned} D_{\text{tie}}(a) &= \frac{3}{4}(t_{e_1} - t_{e_2})^2, \\ D_{\text{adj}}(a) &= D_{\text{tie}}(a) + D_{\text{strength}}(a) + D_{\text{flip}}(a), \quad D_{\text{strength}}(a) = \frac{1}{4}(|b_{e_1}| - |b_{e_2}|)^2, \\ D_{\text{flip}}(a) &= |b_{e_1} b_{e_2}| \mathbb{I}\{b_{e_1} b_{e_2} < 0\}. \end{aligned} \quad (5)$$

For finite measurement, the reported components use a cross-product of independent 15-seed halves while retaining both measurement replicates. The empirical report is therefore $D_{\text{tie}}^{\text{CF}}$ plus a directional-bias component; the nonlinear strength/flip split is retained as population algebra, not presented as an unbiased finite-sample decomposition. The canonical per-seed U-corrected D_{adj} is reconstructed from the merged member-level risks. Finite-measurement component calibration and the comparison between cross-fit and canonical estimates are reported in the Supplement; they are not additional population identities.

Estimand properties. Because each $\pi_{a,e}$ is a probability vector, $0 \leq D_{\text{adj}} \leq 1$. The discrepancy is symmetric in the two contexts and equals zero if and only if the pair-state distributions agree almost everywhere over the declared pair universe. It is therefore a squared divergence; we do not claim a triangle inequality and use “discrepancy” rather than “distance”. The bound applies to the population estimand. Its unbiased finite-sample U-corrected estimate $\hat{D}_{\text{adj}}$ is signed and is not clipped to $[0, 1]$; a negative estimate can occur when the observed pair-state discrepancy is smaller than the finite-replicate correction. In the deterministic tie-free limit, each pair law collapses to a point mass and D_{adj} reduces exactly to the normalized Kendall inversion fraction, $(1 - \tau)/2$.

More generally, it equals cross-context pair disagreement minus one half of each within-context pair disagreement, a population quantity that deterministic Kendall correlation does not define under repeated noisy measurement. Its distinct role is finite-measurement transport, where pair states are uncertain and the matched within-context component must be estimated rather than ignored. This distinction is not cosmetic. Consider one perturbation pair whose strict ordering probabilities are $(0.9, 0, 0.1)$ in the source and $(0.6, 0, 0.4)$ in the target. Both contexts have the same modal strict order, so Kendall after collapsing each pair law to its modal order records no inversion, whereas $D_{\text{adj}}(a) = \frac{1}{2}[(0.9 - 0.6)^2 + (0.1 - 0.4)^2] = 0.09$. The estimand therefore detects transport of ordering reliability even when the modal order does not flip. Conversely, an expression-space MMD can be positive when every risk ordering is unchanged; it tests a different population object. Hence D_{adj} is neither renamed Kendall nor a generic distribution-shift score: it is the declared average separation of finite-measurement pair-state laws. For finite resampling sets, the within-context term uses the unbiased order-2 U-statistic (Hoeffding, 1948). A strict direction is declared stable when the 95% equal-tail Beta(1,1) posterior credible interval excludes 0.5, with at least eight non-tie Monte Carlo resampling outcomes supporting the strict comparison. This threshold diagnoses state-estimation stability and is not a count of biological replicates; otherwise the pair remains unresolved. This categorical state rule is separate from the 90% intervals reported for effect estimates. At the no-transport boundary the U-statistic is degenerate, so an ordinary percentile bootstrap is not used as a zero-effect test. Detectability is instead assessed from the empirical variance of the 30 independent seed-level U-corrected contributions using a Student- t interval. A complete-seed block bootstrap is retained as a sensitivity analysis. This preserves the registered seed as the independent unit and does not resample the two context arms independently.

The primary Frangieh surface contains 243 shared perturbation labels across three contexts. Fig. 1B selects source-rank quantile exemplars only to keep the complete braid legible, whereas its background ribbons and the Fig. 1 source table retain the full 243-label source-frozen surface. Fig. 1D is a compact decision preview; the detailed shortlist flow appears only in Fig. 4A. Aggregate transport estimates use their declared full eligible tables. The predictor uses a ten-component PCA basis, ridge penalty $\alpha = 0.1$, one global five-fold perturbation split and three independent model members. Matched raw-cell pseudoreplicates use 30 resampling seeds and a minimum of 20 cells per perturbation-context. The independent Nadig analysis uses the exact shared HepG2-Jurkat perturbation set and is not pooled with Frangieh.

2.1 DECISION CONSEQUENCES

For this descriptive decision analysis, $\bar{R}_{p,e}^{(m)}$ is the mean full-depth risk over the frozen model members and measurement replicates for metric m . We set $k = \max(1, \lfloor q|\mathcal{P}| \rfloor)$ for budget fraction q and break exact boundary ties by the fixed perturbation identifier. For a source context e_s and selection budget k , let

$$S_k(e_s) = \arg \min_{S: |S|=k} \sum_{p \in S} \bar{R}_{p,e_s}^{(m)} \quad (6)$$

be the source-ranked shortlist. Its target-context loss is $L_{e_t}(S) = k^{-1} \sum_{p \in S} \bar{R}_{p,e_t}^{(m)}$, and normalized selection regret compares $L_{e_t}(S_k(e_s))$ with the target-oracle shortlist $S_k(e_t)$. Thus the decision layer measures the consequence of transporting a source reliability order, rather than evaluating a second predictor.

For the decision layer, the target regret at budget k is defined as

$$\text{Regret}_k(e_s \rightarrow e_t) = L_{e_t}(S_k(e_s)) - L_{e_t}(S_k(e_t)), \quad (7)$$

where the second term is the target-oracle loss. Let $S_k^{\text{worst}}(e_t)$ be the deterministic k -item set with largest target risk and let $L_{e_t}(\mathcal{U})$ be the mean target risk over the eligible universe. We report two explicitly distinct normalizations,

$$\text{Regret}_k^{\text{rand}} = \frac{\text{Regret}_k}{L_{e_t}(\mathcal{U}) - L_{e_t}(S_k(e_t))}, \quad \text{Regret}_k^{\text{worst}} = \frac{\text{Regret}_k}{L_{e_t}(S_k^{\text{worst}}) - L_{e_t}(S_k(e_t))}. \quad (8)$$

The first preserves the random-reference scale used in the primary result; the second is bounded by the deterministic worst shortlist. Rows with a non-positive normalization denominator are undefined and excluded; no such row occurred in the reported surface. Both target-oracle quantities are descriptive benchmarks computed from the same target table, not prospective policies.

Fig. 2 Context-dependent reliability transport and independent replication




Figure 2: Context-dependent reliability transport. (A) Six directed Frangieh transfers with three metric estimates and 90% intervals. (B) Stable-pair inversion fraction: fraction of pairs stable in both contexts whose strict ordering reverses. (C) Same-context independent pseudoreplicate control, shown as point estimates because it is a floor diagnostic rather than a biological transfer. (D) Independent Nadig HepG2–Jurkat replication in both source directions. Metric color is fixed; in panel D, marker shape identifies the source direction and context is otherwise carried by position and labels.

3 RESULTS

3.1 RELIABILITY ORDERING CHANGES ACROSS BIOLOGICAL CONTEXTS

Figure 1 introduces the object of interest: predictions are held fixed while target biology changes, so movement in the perturbation order cannot be attributed to refitting the predictor. The full Frangieh atlas contains 243 shared labels; Fig. 1B retains all 243 background trajectories and highlights a deterministic set of source-rank quantile exemplars so that individual ribbons remain readable. Preprocessing and replication eligibility are summarized in Appendix A; the complete candidate registry is retained with the repository artifacts. Fig. 1D previews the fixed top-10% shortlist replacement example; the complete label-level risk accounting is retained in the figure source table. The same-context resampling control gives $D_{\text{adj}} = 0.0012 - 0.0054$ across its nine context–metric cells, whereas the matched biological transfers produce larger measurement-adjusted values.

The full-size Ctrl-to-IFN γ transfer illustrates metric dependence. For delta cosine, $D_{\text{cross}} = 0.324175$ and $D_{\text{adj}} = 0.093331$ ($[0.056748, 0.134389]$); for Systema centroid, the corresponding values are 0.110196 and 0.035181 ($[0.017110, 0.057823]$); for absolute-effect rank they are 0.328296 and 0.073300 ($[0.035638, 0.111859]$). These generated values are read from the same prespecified analysis table used by the figure builder. The conservative joint-floor stress test is retained separately, because it is not an additive component of the measurement-adjusted estimand. The atlas in Fig. 2 retains metric and transfer heterogeneity rather than averaging it away. The finite-sample estimator and tie rule are given in Appendix B.

3.2 CALIBRATED INFERENCE DISTINGUISHES TRANSPORT FROM FINITE MEASUREMENT

The estimator addresses a finite-measurement target, so a useful comparison is not only against raw rank distances but also against their strongest simple matched-null variants. In a registered 120-condition known-truth simulation, raw Kendall, Spearman and top-10 Jaccard declared transport in 0.952, 0.995 and 1.000 of null conditions, respectively. Label bootstrapping alone did not remove this finite-sample effect. Subtracting a matched within-context null reduced the Kendall false-transport rate to 0.003 and the Jaccard rate to 0.000, while the U-corrected D_{adj} rate was

0.017. At non-zero inversion fractions, detection power was 1.000 for D_{adj} and 0.997 for matched-null Kendall. Thus the result is not that D_{adj} universally dominates every rank statistic: precise measurements make deterministic rank comparators reasonable, whereas D_{adj} supplies one coherent pair-state population estimand with an explicit finite-replicate correction. For the interval itself, direct variance estimation over the 30 independent seed-level U-corrected contributions yielded a null false-transport rate of 0.040 with 0.960 coverage in 2,880 known-truth trials, and 0.015 with 0.985 coverage in 3,000 Frangieh semi-synthetic null trials. These intervals are mildly conservative and replace the anti-conservative hierarchical percentile intervals in every detectability result. The full calibration table, including coverage, bias and the Jaccard comparison, is provided in the Supplement.

3.3 INDEPENDENT RAW-CELL SYSTEMS REPRODUCE ORDERING TRANSPORT

The independent Nadig replication provides a different biological system and context pair and a larger shared perturbation universe. The adjusted discrepancy D_{adj} is 0.176 for delta cosine, 0.138 for Systema centroid and 0.129 for absolute-effect rank, with positive 90% intervals in both source directions. This replication supports the ordering-level phenomenon without being used to tune the Frangieh decision association. Matched-universe sensitivity for this replication is reported in Appendix C, while the complete two-direction full-size surface remains in the figure source table.

A second raw-cell replication reconstructs the complete hierarchy in the Tian iPSC and day-7 neuron CRISPR screens (Tian et al., 2019), rather than using an aggregate rank audit. We froze 26 shared single-gene perturbations, a 1,536-gene panel chosen from source controls only, and a 40-cell budget before evaluating either direction. Five of six direction-metric intervals were positive. From iPSC to neuron, D_{adj} was 0.032 [0.005, 0.060], 0.010 [−0.010, 0.030], and 0.016 [0.004, 0.028] for delta cosine, Systema and absolute-effect rank, respectively; the reverse direction gave 0.049 [0.018, 0.079], 0.042 [0.025, 0.060], and 0.013 [0.001, 0.025]. The direction-specific estimates support context-dependent ordering transport; their asymmetry is reported descriptively rather than treated as a separate biological conclusion. The complete six-row ledger and source-frozen construction are in Appendix C.

3.4 SOURCE-FROZEN PREDICTOR FAMILIES PRESERVE THE QUALITATIVE FINDING

As one matched predictor-family sensitivity, we also evaluated a source-only RBF kernel-ridge family on the same Frangieh surface. Its source-only penalty selection required an expanded grid. The matched transport profile had Spearman $\rho = 0.983$ with the primary bilinear-ridge profile (within-metric profile correlations 0.829–1.000; mean row-wise prediction-vector cosine 0.733 in the Supplement), and adjusted-discrepancy point and interval signs agreed on all 18 transfer-metric rows. This is descriptive evidence that the profile is not an immediate artifact of one parameterization, not predictor- family equivalence or broader generalization.

We then ran the official GEARS graph predictor under the same strict information contract. For each Frangieh source context, five label-disjoint folds and all graph construction, fitting and tuning used source cells only; 13 labels absent from the official interaction graph were excluded before splitting, leaving 230 perturbations and 8,229 evaluation genes. The frozen out-of-fold vectors were evaluated by the same 30-seed, two-pseudoreplicate raw-cell bridge. All 18 directed transfer-metric rows retained positive adjusted-discrepancy intervals. Relative to the bilinear profile, GEARS had perfect point-sign agreement but lower magnitude concordance (mean within-metric Spearman $\rho = 0.371$: −0.086, 1.000 and 0.200 for delta cosine, Systema centroid and absolute-effect rank). Thus a modern graph model reproduces the existence of transport while showing that its detailed geometry is predictor- and metric-dependent; it does not establish family equivalence or superiority. The complete surface is reported in Supplementary Appendix H.

3.5 CELL BUDGET DETERMINES WHEN REORDERING IS RESOLVABLE

We next hold the perturbation universe fixed and resample 10, 20, 40, 80 and 160 cells per perturbation-context pseudoreplicate, together with a matched full-depth endpoint. The signed $D_{\text{adj}}(n)$ trajectory is retained. The first budget within ten percent of the full-depth value summarizes resolution; detectability requires a positive lower endpoint of the 90% interval. These are

Fig. 3 Cell budget determines when reordering is resolvable




Figure 3: Cell budget defines a finite-measurement resolution boundary. (A) Signed depth trajectories for the primary Ctrl–IFN γ contrast with other transfers as pale paths. (B) The six-transfer resolution map; circles mark detectability and squares mark the resolution budget, with open glyphs denoting not reached. (C) Cross-context transport and matched same-context split-half pair-law null margins across depth for the five prespecified comparator families. This is a finite-measurement diagnostic, not the canonical per-seed D_{adj} summary. The full-depth endpoint is the fixed-160 eligibility universe ($n = 139$), not the wider full-size minimum-20 universe used for Figs. 1–2. Signed finite-sample estimates are retained, so null diagnostics can be negative. Detectability and quantitative resolution are separate experimental requirements under the declared resampling protocol.

metric- and transfer-specific summaries, not a universal cell threshold. We additionally compare each transport statistic with a matched same-context pseudo-context null at every depth, because the canonical within-context measurement term is not itself a same-context D_{adj} estimate.

The trajectories rise at different rates for the three metrics, and several source–contrast cells remain unresolved at the available budgets. Increasing cell budget therefore improves resolution under this protocol without producing a single portable cutoff. This measurement boundary is a prerequisite for interpreting a context-associated ordering difference beyond the measured sampling floor. The matched-null analysis shows positive cross-minus-null margins for every comparator at every depth, but does not support presenting D_{adj} as a universally superior rank statistic; its role is the explicit tie-aware, measurement-adjusted estimand. The depth comparator audit is reported in the Supplementary finite-measurement audit. The depth definitions and all transfer-specific results are in Appendix D.

3.6 CONTEXT-DEPENDENT REORDERING CHANGES SHORTLIST DECISIONS

To connect the ordering to an experimental action, each source context selects the lowest-risk perturbations at budgets $k/n \in \{0.05, 0.10, 0.20, 0.50\}$ and the target context evaluates the selected set. We report target top- k retention and normalized regret under both a random-reference and a bounded worst-oracle normalization.

The primary full-depth link contains six directed source–target surfaces and three metrics, hence 18 descriptive rows. Across all rows, adjusted reordering and normalized regret have pooled Spearman $\rho = 0.78$; after ranking within metric, the association is $\rho = 0.26$. the three unordered context-pair

Fig. 4 From reordering to experimental regret




Figure 4: Context-dependent ordering changes experimental priorities. (A) Detailed example of delta-cosine top- k replacement under the Ctrl-to-IFN γ shift (6 of 25 shortlist positions are retained; 19 are dropped and 19 are new). (B) The 18 full-depth transfer-metric observations; marker shape identifies metric and pale connectors identify directed transfer. The pointwise error bars are canonical 90% intervals; the descriptive association and its unordered-context-pair sensitivity are reported in the text, and no fitted line is implied. (C) Unordered-context-pair-held-out incremental prediction of normalized regret beyond metric identity; intervals are block-bootstrap diagnostics.

and source-context sensitivities are reported separately. The association is a decision-level summary of the declared surface, not a claim that every predictor or biological system will have the same slope.

An explicit unordered-context-pair leave-one-out comparison separates metric identity, mean-fidelity shift, and ordering transport. Adding the matched mean-risk shift to metric identity reduces OOF MAE by 0.013 ($[-0.001, 0.027]$). Adding D_{adj} increased OOF MAE from 0.087 to 0.092, corresponding to an improvement score of -0.005 ($[-0.014, 0.003]$); the two nested contrasts are shown separately in Fig. 4. Thus the pooled association is not interpreted as a unique causal contribution of one scalar ordering summary. The corresponding unordered-pair block-bootstrap table is reported in the Supplement; no headline permutation p -value is used for only three blocks.

The shortlist view makes the mechanism concrete: perturbations that occupy the source top- k need not occupy the target top- k , so a source ranking can lose retention and incur regret even when the predictor itself has not changed. The budget-wise association summaries show that the link is not confined to one selection fraction, while its magnitude depends on the transfer and metric. The four-budget, two-normalization robustness summary is reported in Appendix D. A metric-matched mean-risk control, using the same eligible labels and frozen predictions, shows that context-average fidelity and reliability- order transport are empirically non-equivalent (Supplementary Fig. S3).

Additional target-informed and prospective source-only analyses are reported in the Supplement. They are kept separate from the post-hoc transport claim: the target-informed audit uses target outcomes, whereas the prospective analysis uses source-observable features only and reserves target outcomes exclusively for evaluation.

4 DISCUSSION

The central result is that perturbation reliability is an environment-dependent ordering, not a fixed property of a perturbation. With predictions held fixed, matched biological context shifts reordered realized risk beyond a same-context sampling floor, and the direction recurs in both directions of an independent HepG2–Jurkat context pair. The result is therefore an empirical reliability-transport object: it does not claim that every predictor or biological system must reorder in the same way.

Unlike response-prediction benchmarks (Lotfollahi et al., 2019; 2023; Bunne et al., 2023; Roohani et al., 2024; Cui et al., 2024; Wei et al., 2026), our object asks whether perturbation-level error ordering remains useful after biology changes. It is also distinct from experimental perturbation reproducibility (Wang et al., 2026): predictor, metric and context are part of the estimand.

Cell budget separates detectability from resolution of transport magnitude. The tie-aware support rule and matched same-context control make this distinction explicit. In the deterministic tie-free limit, D_{adj} reduces to the normalized Kendall inversion fraction; here it extends the same question to uncertain pair states under finite repeated measurement.

Source-ranked top- k choices lose target retention when reliability moves, but this is a diagnostic rather than a decision-focused learner (Donti et al., 2017; Wilder et al., 2019; Elmachtoub & Grigas, 2022). The 18 rows arise from only three unordered context-pair blocks. Metric-stratified associations are weaker than the pooled surface; mean-risk shift provides the only modest OOF improvement (0.013; $[-0.001, 0.027]$) and the additional D_{adj} contrast beyond mean shift does not add held-out value (-0.005 ; $[-0.014, 0.003]$). We therefore do not claim that one scalar universally predicts regret.

Scope and testable extensions. The primary decision surface is limited to one bilinear-ridge predictor, three metrics and three unordered Frangieh context pairs. Matched RBF, neural and GEARS analyses test predictor sensitivity but are not pooled with it; the Nadig and Tian reconstructions replicate biological transport but do not enlarge the decision sample. TxPert lacked a source-matched Frangieh checkpoint, and the outcome-blind Feng aggregate audit was neutral, so neither is promoted as a positive replication. The next test is a prospectively frozen, label-disjoint analysis across additional contexts and modalities. This scope is important as context diversity becomes central to virtual-cell evaluation (Dibaeinia et al., 2026).

Reproducibility. The anonymous core release provides fixed splits, source tables and commands for the four main figures, Supplementary Figs. S1–S7 and compact tables. Appendices A, B and I specify the data, estimators and regeneration path; engineering artifacts are excluded from manuscript evidence.

Code and data availability. The same anonymous release contains analysis and figure code, manuscript sources and compact result products. Raw matrices are not redistributed. Frangieh and Nadig data are available from the public scPerturb v1.4 Zenodo record; the repository record; Feng data are available from its Cell Genomics record (Feng et al., 2026). The repository README lists exact inputs and commands.

AI use. AI-assisted tools supported literature triage, mathematical and experimental brainstorming, implementation and debugging of analysis code, figure-script development, interpretation checks, and language editing. The authors designed the study, inspected all generated analyses, verified citations and numerical claims against source artifacts, and take responsibility for the manuscript.

REFERENCES

Britt Adamson, Thomas M. Norman, Marco Jost, Min Y. Cho, James K. Nunez, Yuwen Chen, Jacqueline E. Villalta, Luke A. Gilbert, Max A. Horlbeck, Marco Y. Hein, Richard A. Pak, Andrew N. Gray, Carol A. Gross, Atray Dixit, Oren Parnas, Aviv Regev, and Jonathan S. Weissman. A multiplexed single-cell CRISPR screening platform enables systematic dissection of the unfolded protein response. *Cell*, 167(7):1867–1882.e21, 2016. doi: 10.1016/j.cell.2016.11.048.

486 Constantin Ahlmann-Eltze, Wolfgang Huber, and Simon Anders. Deep-learning-based gene per-
487 turbation effect prediction does not yet outperform simple linear baselines. *Nature Methods*, 22:
488 1657–1661, 2025. doi: 10.1038/s41592-025-02772-6.

489

490 Charlotte Bunne, Stefan G. Stark, Gabriele Gut, Jacobo Sarabia del Castillo, Mitch Levesque,
491 Kjong-Van Lehmann, Lucas Pelkmans, Andreas Krause, and Gunnar Raetsch. Learning single-
492 cell perturbation responses using neural optimal transport. *Nature Methods*, 20(11):1759–1768,
493 2023. doi: 10.1038/s41592-023-01969-x.

494

495 Haotian Cui, Chloe Wang, Hassaan Maan, Kuan Pang, Fengning Luo, Nan Duan, and Bo Wang.
496 scGPT: toward building a foundation model for single-cell multi-omics using generative AI. *Nature Methods*, 21(8):1470–1480, 2024. doi: 10.1038/s41592-024-02201-0.

497

498 Payam Dibaeinia, Sudarshan Babu, Mei Knudson, Ali ElSheikh, Yibo Wen, Han Liu, Jason Perera,
499 and Aly A. Khan. Virtual cells need context, not just scale. *bioRxiv*, 2026. doi: 10.64898/2026.
500 02.04.703804.

501

502 Atray Dixit, Oren Parnas, Biyu Li, Jenny Chen, Charles P. Fulco, Livnat Jerby-Arnon, Nemanja D.
503 Marjanovic, Danielle Dionne, Tyler Burks, Raktima Raychowdhury, Britt Adamson, Thomas M.
504 Norman, Eric S. Lander, Jonathan S. Weissman, Nir Friedman, and Aviv Regev. Perturb-Seq:
505 Dissecting molecular circuits with scalable single-cell RNA profiling of pooled genetic screens.
506 *Cell*, 167(7):1853–1866.e17, 2016. doi: 10.1016/j.cell.2016.11.038.

507

508 Priya L. Donti, Brandon Amos, and J. Zico Kolter. Task-based end-to-end model learning
509 in stochastic optimization. In *Advances in Neural Information Processing Systems*, vol-
510 ume 30, pp. 5484–5494, 2017. URL [https://papers.nips.cc/paper/2017/hash/
511 3fc2c60b5782f641f76bcefc39fb2392-Abstract.html](https://papers.nips.cc/paper/2017/hash/3fc2c60b5782f641f76bcefc39fb2392-Abstract.html).

512

513 Adam N. Elmachtoub and Paul Grigas. Smart “predict, then optimize”. *Management Science*, 68
514 (1):9–26, 2022. doi: 10.1287/mnsc.2020.3922.

515

516 Claudia Feng, Elin Madli Peets, Yan Zhou, et al. A genome-scale single-cell CRISPRi map of
517 trans gene regulation across human pluripotent stem cell lines. *Cell Genomics*, 6(2):101076,
518 2026. doi: 10.1016/j.xgen.2025.101076. URL [https://pmc.ncbi.nlm.nih.gov/
519 articles/PMC12903452/](https://pmc.ncbi.nlm.nih.gov/articles/PMC12903452/). Public multi-line, multi-donor targeted-screen resource accessed
520 for an outcome-blind external replication audit.

521

522 Chris J. Frangieh, Johannes C. Melms, Pratiksha I. Thakore, Kathryn R. Geiger-Schuller, Patricia
523 Ho, Adrienne M. Luoma, Brian Cleary, Livnat Jerby-Arnon, Shruti Malu, Michael S. Cuoco,
524 Maryann Zhao, Casey R. Ager, Meri Rogava, Lila Hovey, Asaf Rotem, Chantale Bernatchez,
525 Kai W. Wucherpennig, Bruce E. Johnson, Orit Rozenblatt-Rosen, Dirk Schadendorf, Aviv
526 Regev, and Benjamin Izar. Multimodal pooled Perturb-CITE-seq screens in patient models
527 define mechanisms of cancer immune evasion. *Nature Genetics*, 53:332–341, 2021. doi:
528 10.1038/s41588-021-00779-1.

529

530 Arthur Gretton, Karsten M. Borgwardt, Malte J. Rasch, Bernhard Schölkopf, and Alexander Smola.
531 A kernel two-sample test. *Journal of Machine Learning Research*, 13:723–773, 2012.

532

533 Wassily Hoeffding. A class of statistics with asymptotically normal distribution. *The Annals of*
534 *Mathematical Statistics*, 19(3):293–325, 1948. doi: 10.1214/aoms/1177730196.

535

536 Mohammad Lotfollahi, F. Alexander Wolf, and Fabian J. Theis. scGen predicts single-cell pertur-
537 bation responses. *Nature Methods*, 16:715–721, 2019. doi: 10.1038/s41592-019-0494-8.

538

539 Mohammad Lotfollahi, Anna Klimovskaia Susmelj, Carlo De Donno, Leon Hetzel, Yuge Ji, Igna-
540 cio L. Ibarra, Sanjay R. Srivatsan, Mohsen Naghipourfar, Riza M. Daza, Beth Martin, Jay Shen-
541 dure, Jose L. McFaline-Figueroa, Pierre Boyeau, F. Alexander Wolf, Nafissa Yakubova, Stephan
542 Ginnemann, Cole Trapnell, David Lopez-Paz, and Fabian J. Theis. Predicting cellular responses
543 to complex perturbations in high-throughput screens. *Molecular Systems Biology*, 19(6):e11517,
544 2023. doi: 10.15252/msb.202211517.

540 Ajay Nadig, Joseph M. Replogle, Angela N. Pogson, Mukundh Murthy, Steven A. McCarroll,
541 Jonathan S. Weissman, Elise B. Robinson, and Luke J. O'Connor. Transcriptome-wide analy-
542 sis of differential expression in perturbation atlases. *Nature Genetics*, 57:1228–1237, 2025. doi:
543 10.1038/s41588-025-02169-3.

544 Stefan Peidli, Tessa D. Green, and Ciyue Shen. scPerturb: harmonized single-cell perturbation data.
545 *Nature Methods*, 21(3):531–540, 2024. doi: 10.1038/s41592-023-02144-y.

546 Joseph M. Replogle, Reuben A. Saunders, Angela N. Pogson, Jeffrey A. Hussmann, Alexander
547 Lenail, Alina Guna, Lauren Mascibroda, Eric J. Wagner, Karen Adelman, Gila Lithwick-Yanai,
548 Nika Iremadze, Florian Oberstrass, Doron Lipson, Jessica L. Bonnar, Marco Jost, Thomas M.
549 Norman, and Jonathan S. Weissman. Mapping information-rich genotype-phenotype landscapes
550 with genome-scale Perturb-seq. *Cell*, 185(14):2559–2575.e28, 2022. doi: 10.1016/j.cell.2022.
551 05.013.

552 Yusuf Roohani, Kexin Huang, and Jure Leskovec. Predicting transcriptional outcomes of novel
553 multigene perturbations with GEARS. *Nature Biotechnology*, 42(6):927–935, 2024. doi: 10.
554 1038/s41587-023-01905-6.

555 Aravind Sankaran, Lars Karlsson, and Paolo Bientinesi. Ranking with ties based on noisy perfor-
556 mance data. *International Journal of Data Science and Analytics*, 20:4363–4384, 2025. doi:
557 10.1007/s41060-025-00722-1.

558 Jordan W. Squair, Michael A. Skinnider, Matthieu Gautier, Leonard J. Foster, and Grégoire Courtine.
559 Prioritization of cell types responsive to biological perturbations in single-cell data with Augur.
560 *Nature Protocols*, 16(8):3836–3873, 2021. doi: 10.1038/s41596-021-00561-x.

561 Ruilin Tian, Mariam A. Gachechiladze, Connor H. Ludwig, Matthew T. Laurie, Jason Y. Hong,
562 Diane Nathaniel, Anika V. Prabhu, Michael S. Fernandopulle, Rajan Patel, Mehrnoosh Abshari,
563 Michael E. Ward, and Martin Kampmann. Crispr interference-based platform for multimodal
564 genetic screens in human ipsc-derived neurons. *Neuron*, 104(2):239–255.e12, 2019. doi: 10.
565 1016/j.neuron.2019.07.014.

566 Ramon Vinas Torne, Maciej Wiatrak, Zoe Piran, Shuyang Fan, Liangze Jiang, Sarah A. Teichmann,
567 Mor Nitzan, and Maria Brbic. Systema: A framework for evaluating genetic perturbation response
568 prediction beyond systematic variation. *Nature Biotechnology*, 44:1050–1059, 2026. doi: 10.
569 1038/s41587-025-02777-8.

570 Xi Wang, Jack Kuipers, Florian Hugi, Randall Jeffrey Platt, and Niko Beerenwinkel. Reliable single-
571 cell perturbations explain and improve model performance. *bioRxiv*, 2026. doi: 10.64898/2026.
572 08.11.744177.

573 Zhiting Wei, Yiheng Wang, Yicheng Gao, Shuguang Wang, Ping Li, et al. Benchmarking algorithms
574 for generalizable single-cell perturbation response prediction. *Nature Methods*, 23:451–464,
575 2026. doi: 10.1038/s41592-025-02980-0. URL <https://www.nature.com/articles/s41592-025-02980-0>.

576 Frederik Wenkel, Wilson Tu, Cassandra Masschelein, Hamed Shirzad, Liam Hodgson, Ihab Bendidi,
577 Cian Eastwood, Shawn T. Whitfield, Craig Russell, Yassir El Mesbahi, Jiarui Ding, Marta M.
578 Fay, Berton Earnshaw, Emmanuel Noutahi, and Alisandra K. Denton. TxPert: using multiple
579 knowledge graphs for prediction of transcriptomic perturbation effects. *Nature Biotechnology*,
580 2026. doi: 10.1038/s41587-026-03113-4.

581 Bryan Wilder, Bistra Dilkina, and Milind Tambe. Melding the data-decisions pipeline: Decision-
582 focused learning for combinatorial optimization. In *Proceedings of the AAAI Conference on*
583 *Artificial Intelligence*, volume 33, pp. 1658–1665, 2019. doi: 10.1609/aaai.v33i01.33011658.

584 Daniel R. Wong, Abby S. Hill, and Rob Moccia. Simple controls exceed best deep learning algo-
585 rithms and reveal foundation model effectiveness for predicting genetic perturbations. *Bioinform-*
586 *atics*, 41(6):btaf317, 2025. doi: 10.1093/bioinformatics/btaf317.

587 Or Zuk, Liat Ein-Dor, and Eytan Domany. Ranking under uncertainty. In *Proceedings of the 23rd*
588 *Conference on Uncertainty in Artificial Intelligence*, pp. 466–473, 2007.